



BY FENIL GAJJAR

# KUBERNETES DAILY TASKS

## KUBERNETES INGRESS

- COMPREHENSIVE GUIDE
- THEORY + PRACTICAL TASKS
- REAL TIME SCENARIO TASKS
- FOLLOW FOR MORE TASKS
- [linkedin.com/in/fenil-gajjar](https://www.linkedin.com/in/fenil-gajjar)



CONTACT US

[fenilgajjar.devops@gmail.com](mailto:fenilgajjar.devops@gmail.com)



# Kubernetes Ingress

**Unlocked:**

**Seamless Traffic Routing**

**Like a Pro!** 

---

# 🌟 Welcome to Another Milestone in Our Kubernetes Journey! 🚀

First, I want to express my **deep gratitude** for all the support and engagement you've shown throughout this Kubernetes series. 🙌 Your enthusiasm and encouragement drive me to bring even more valuable insights every day. This journey wouldn't be the same without your support, so **thank you for being part of it!**

## 🌍 Today's Focus: Kubernetes Ingress

As we continue mastering Kubernetes, today we're unlocking a **powerful networking component—Ingress!** 💡 In the real world, applications need a way to be **accessible from the outside world**, whether through **HTTP, HTTPS, or custom routes**. Kubernetes provides multiple ways to expose services, but Ingress takes it **a step further**, offering a more **efficient, flexible, and centralized approach** for managing external access.

---

In this guide, we'll dive deep into:

- ✅ What Kubernetes Ingress is and why it's essential
- ✅ How it works behind the scenes
- ✅ Key configurations and best practices
- ✅ Hands-on examples to set up and manage Ingress effectively

Whether you're a beginner or already familiar with Kubernetes networking, this will **level up your understanding and skills**. So, get ready to explore, experiment, and enhance your Kubernetes expertise! 🔥

Stay tuned, and let's keep growing together. 💪 **Keep supporting, keep learning, and let's master Kubernetes one step at a time!** 🚀



# Why Do We Need Ingress?

Before Kubernetes Ingress came into play, exposing applications to the outside world was **not so easy**. Let's take a quick trip back in time and see how things worked before Ingress and why it became a game-changer! 🎯



## The Pre-Ingress Era: How Did We Expose Services

### Before?

In Kubernetes, every Pod has its **own private IP**. But here's the problem:



**Pods are ephemeral**—they come and go!



Their IPs **keep changing**, making direct access impossible.

So, Kubernetes introduced **Services** to solve this issue by providing a **stable way** to connect to Pods. But how do we expose these services outside the cluster?



### Option 1: NodePort

Kubernetes allowed us to expose services using **NodePort**, which opens a fixed port on each node. The issue?



The port range is limited (30000-32767).

---

✗ It's **not user-friendly**—imagine remembering multiple ports for different services!

✗ Not scalable for production.

## Option 2: LoadBalancer

Next, we had **LoadBalancer Services**, which brought us a step closer to real-world usage. A cloud provider assigns an **external IP**, and traffic is routed to the service.

But still...

✗ Every service needs a **separate LoadBalancer**.

✗ **Expensive**—creating a LoadBalancer for each service is costly.

✗ No built-in traffic routing features like **path-based or domain-based routing**.

## Enter Ingress: The Game Changer!

To solve these challenges, **Kubernetes Ingress** was introduced! 

✨ Instead of exposing each service separately, Ingress acts as a **smart gateway** that routes external traffic **inside the cluster** based on:

✓ **Hostnames** (e.g., **app.example.com**)

✓ **Paths** (e.g., **/api**, **/dashboard**)

---

✓ **TLS (HTTPS) termination**

✓ **Load balancing** between multiple services

Think of **Ingress as the front door** to your applications—it lets you define flexible, user-friendly routing rules in a **single place**.

So now, instead of managing multiple LoadBalancers or dealing with NodePorts, we just set up an **Ingress Controller** and define simple rules. **Boom!** 🎯

🚀 **Less cost, less complexity, more power!**



# What Exactly is Ingress in Kubernetes?

In Kubernetes, **Ingress** is a powerful API object that manages how external users access services within a cluster. It acts as a **smart gateway** that routes traffic based on **hostnames, paths, and rules**—eliminating the need for multiple LoadBalancers and messy NodePort configurations.



## Understanding Ingress in Simple Terms

Think of **Ingress** as the receptionist in a large office building. When visitors arrive, the receptionist:

- ✓ **Greets them** (accepts incoming traffic)
- ✓ **Checks their purpose** (examines the requested hostname or path)
- ✓ **Directs them to the right department** (routes traffic to the correct service)

Without Ingress, visitors would have to **directly enter individual rooms** (NodePort) or **use separate entrances** (multiple LoadBalancers), which is inefficient.



---

## Key Features of Ingress (Made Simple & Fun!)

### 1 Smart Traffic Routing


Imagine a **mall entrance** with different signs leading visitors to the right store.

Ingress works the same way by directing users to the correct service based on:

✓ **Paths** (/shop → frontend, /api → backend)

✓ **Hostnames** (shop.myapp.com → shop service, api.myapp.com → API service)

### 2 Built-in Security with HTTPS (SSL/TLS Termination)

Ever noticed the  lock icon in your browser? That's **HTTPS encryption**, keeping data safe. Ingress makes it easy to handle SSL certificates, ensuring secure communication without configuring it in each service.

### 3 Load Balancing for Performance

Just like a traffic cop managing cars at a busy intersection, Ingress **distributes** traffic evenly across multiple instances of a service. This prevents overload and ensures **smooth performance**.

---

## 4 Single External IP (Saves Cost & Complexity)

Without Ingress: Every service needs a **separate LoadBalancer**, increasing cloud costs.

With Ingress: One **central LoadBalancer** handles all traffic, making it **cost-effective & scalable**.

## 5 Easy Traffic Control (Rewrites, Redirects, & More!)

Ingress lets you:

 **Redirect users** (e.g., [http](#) → [https](#))

 **Rewrite URLs** ([/old-page](#) → [/new-page](#))

 **Rate-limit requests** to prevent overload

## 6 Works with Ingress Controllers

Ingress itself is just a set of rules, but it needs an **Ingress Controller** to enforce them. Popular controllers include:

-  **NGINX Ingress Controller** (Most common)

- 
- 🚀 **Traefik** (Lightweight & dynamic)
  - 🔥 **Istio Gateway** (For advanced networking)

## 🎯 Why Use Ingress? The Ultimate Traffic Manager in Kubernetes!

Imagine running a **bustling online store** with different sections like the **shop**, **API**, and **admin dashboard**. Without **Ingress**, every section would need its own separate **door** (LoadBalancer or NodePort), making things **messy, expensive, and hard to manage**.

But with **Ingress**, you get **one grand entrance** that intelligently routes visitors to the right place—saving money, improving security, and keeping traffic under control. 🚀

### 💡 Key Advantages of Using Ingress (The Problem It Solves!)

🏠 1 **Single Entry Point Instead of Multiple Exposures**

---

🚫 **Without Ingress:** Each service needs its own LoadBalancer or NodePort, making it expensive and difficult to manage.

✅ **With Ingress:** A **single entry point** smartly routes traffic inside the cluster. No need for multiple external IPs!

**Example:**

- `shop.example.com` → Routes traffic to the frontend service
- `api.example.com` → Routes traffic to the backend service
- `admin.example.com` → Routes traffic to the admin panel

🔒 2 **Secure Communication with HTTPS & SSL/TLS Termination**


🚫 **Without Ingress:** Each service needs separate SSL certificates, making security hard to manage.


✅ **With Ingress:** You can manage **SSL/TLS certificates centrally**, ensuring all traffic is **secure & encrypted**.

**Think of it like this:** Instead of every store in a mall installing their own security system, the mall provides a **centralized security gate** for all!

---


### 3 Load Balancing for Even Traffic Distribution


 **Without Ingress:** Traffic might overwhelm certain services while others stay idle.

 **With Ingress:** Load balancing ensures requests are **distributed evenly**, preventing overload and boosting performance.

**Example:** If 1,000 users hit your application, Ingress ensures they are routed across multiple backend servers efficiently.

### 4 Intelligent Traffic Control (Path-Based & Host-Based Routing)

 **Without Ingress:** You need separate configurations for each service's traffic flow.

 **With Ingress:** You define **routing rules** in one place for organized traffic management.


**Example:**

- `example.com/shop` → Goes to the **frontend**
- `example.com/api` → Goes to the **backend**
- `example.com/admin` → Goes to the **admin dashboard**

---

This keeps things **structured, simple, and easy to scale!**

## 5 **Cost-Effective (No More Expensive LoadBalancers!)**


 **Without Ingress:** You need a LoadBalancer for each exposed service, increasing cloud costs.

 **With Ingress:** A **single** LoadBalancer handles all services, drastically reducing expenses.

**Imagine paying for multiple toll gates vs. having one central highway entrance for all roads!**

## 6 **URL Rewriting & Redirects (Better User Experience!)**

 **Without Ingress:** If a URL structure changes, users might get broken links.

 **With Ingress:** You can set up **redirects & rewrites** to ensure smooth navigation.


**Example:**


- Redirect <http://example.com> → <https://example.com>



- 
- Rewrite `/old-page` → `/new-page` without breaking links

## 7 Traffic Splitting & Canary Deployments

 **Without Ingress:** Rolling out new features requires downtime or complex setups.

 **With Ingress:** You can **split traffic** between different versions of your app (e.g., 80% to v1, 20% to v2) for safe deployments.

**Perfect for testing new features without impacting all users!**

## 8 Works with Ingress Controllers for Extra Power!

Ingress itself is just a set of rules, but it works with **Ingress Controllers** to enforce them:

 **NGINX Ingress Controller** – The most popular & widely used

 **Traefik** – Lightweight & dynamic routing

 **Istio Gateway** – Advanced traffic management for microservices

These controllers enhance **performance, security, and scalability!**

---

## Ingress Architecture – Key Components

### 1 Ingress Controller

- The brain behind **Ingress** that processes incoming traffic and applies routing rules.
- Common controllers: **NGINX**, **Traefik**, **HAProxy**, **AWS ALB Controller**.

### 2 Ingress Resource

- A **YAML configuration file** that defines how traffic should be routed.
- Example: Redirect **app.example.com** to **Service A** and **shop.example.com** to **Service B**.

### 3 Ingress Class

- Helps define **which Ingress Controller** should handle a specific Ingress resource.
- Example: Use **NGINX** for general apps and **AWS ALB** for cloud-based services.

### 4 Service

- Connects the **Ingress** to actual backend applications running in **Pods**.
- Works as a bridge between the **Ingress Controller** and applications.



---

## 5 Pod (Application) 🚀

- The final destination where requests are processed.
- Runs the application that serves the actual content to users.

## 🚀 What is an Ingress Controller & Why is it Important?

If **Ingress is the traffic manager** of Kubernetes, then **Ingress Controller is the engine** that makes it all work! 💡

Think of it like this: **Ingress is the road system**, but without traffic lights and signboards, it's chaos. 🚦 That's where **Ingress Controller** comes in – it ensures traffic flows smoothly to the right places!

## 📌 What is an Ingress Controller?

An **Ingress Controller** is a specialized **Kubernetes component** that processes the rules defined in the **Ingress resource** and handles incoming traffic accordingly. 🌐

---

💡 **Without an Ingress Controller, an Ingress resource alone does NOTHING!**

It's like having a train track (Ingress) but **no train operator** (Ingress Controller). 🚂

## 🎯 How Ingress Controller Works (Step-by-Step)

Let's say you have an e-commerce app with different services:

- **Frontend:** `shop.example.com`
- **Backend API:** `api.example.com`

Here's how an Ingress Controller handles this traffic:

- 1 User requests `shop.example.com` 🌐
  - 2 The request **first reaches the Ingress Controller**. 🚦
  - 3 The Controller **checks the Ingress rules**. 📜
  - 4 It **routes traffic** to the correct service (`Frontend Service`). 🏢
  - 5 The **service forwards traffic** to the right **Pods** (your app). 🚀
- ✅ **Boom! Your app works seamlessly!**

---

## 🔥 Why is an Ingress Controller Important?

- 📌 Without it, Kubernetes Ingress won't work!
- 📌 It **efficiently handles and routes external traffic** to internal services.
- 📌 It provides **SSL termination, authentication, and load balancing**.
- 📌 Reduces **costs** by eliminating multiple cloud load balancers.
- 📌 Enables **path-based** and **host-based routing**.

Imagine running **10 microservices** – instead of 10 different LoadBalancers, you can just use **one Ingress Controller** to manage all the traffic! 🚀

## ⚡ Popular Ingress Controllers

- ♦ **NGINX Ingress Controller** – Most widely used, stable & feature-rich.
- ♦ **Traefik** – Lightweight & supports auto-discovery.
- ♦ **HAProxy** – High-performance with advanced traffic control.
- ♦ **AWS ALB Ingress Controller** – For integrating AWS Load Balancer.
- ♦ **Istio Gateway** – Used in service meshes for advanced traffic control.

---

## How Ingress Works in Kubernetes

### ♦ Step 1: User Makes a Request

When a user tries to access a service (e.g., `app.example.com`), their request first reaches the **Kubernetes cluster's external entry point**.

👉 **Without Ingress:** Every service would need a separate LoadBalancer, increasing complexity and cost.

👉 **With Ingress:** One centralized entry point directs traffic efficiently. 🚦

### ♦ Step 2: The Request Hits the Ingress Controller

The request doesn't go straight to the application! Instead, it first reaches the **Ingress Controller**, which is responsible for processing the request. 🚀

♦ **Ingress Controller reads the rules** defined in the **Ingress resource**.

♦ It determines **where to send the traffic** based on the **domain name, path, or custom rules**.

### ✅ Example:

- Requests to `app.example.com` → Forwarded to the **frontend service**
- Requests to `api.example.com` → Sent to the **backend service**

- 
- Requests to `app.example.com/images` → Redirected to an **image storage service**

### ♦ **Step 3: Ingress Controller Routes Traffic**

Once the Ingress Controller knows where to send the request, it **routes the traffic** to the correct **Kubernetes Service**, which then forwards it to the appropriate **Pods** running the application.

### 💡 **Think of Ingress as an Air Traffic Controller!**

It ensures every request lands at the right destination safely and efficiently. ✈️

### ✅ **Path-based routing example:**


- `/home` → Frontend Service
- `/api` → Backend API Service
- `/static` → CDN or Image Service

### ♦ **Step 4: Response is Sent Back to the User**

---

Once the Pod processes the request, the response **follows the same path back**:

 **Pod → Service → Ingress Controller → User**

This ensures **secure, optimized, and seamless communication** between external users and internal applications. 

---

# Complete YAML Configuration for Kubernetes Ingress Resource

Ingress in Kubernetes allows you to define routing rules to direct external traffic to internal services efficiently. Below is a **complete and well-structured YAML configuration** for an **Ingress resource**, assuming you're using **Kops** as your **Kubernetes cluster**.

## Prerequisites Before Applying Ingress

Before you create an Ingress resource, make sure:

- ✓ You have an **Ingress Controller** (such as Nginx Ingress Controller) installed and running in your cluster.
- ✓ You have **services running** that need to be exposed externally.

## Complete Ingress YAML Configuration

The following example defines an **Ingress resource** that:

- 
- Routes traffic to **multiple services** based on the request's domain or path.
  - Supports **HTTPS with TLS termination** using a Kubernetes secret.
  - Handles **path-based and host-based routing** for different microservices.

apiVersion: networking.k8s.io/v1

kind: Ingress

metadata:

name: my-app-ingress

namespace: default

annotations:

nginx.ingress.kubernetes.io/rewrite-target: /

nginx.ingress.kubernetes.io/ssl-redirect: "true"

spec:

ingressClassName: nginx

tls:

- hosts:

- app.example.com



---

```
    secretName: tls-secret # TLS certificate stored in a
Kubernetes Secret
```

```
rules:
```

```
- host: app.example.com
```

```
  http:
```

```
    paths:
```

```
      - path: /          # Root path for frontend
```

```
        pathType: Prefix
```

```
        backend:
```

```
          service:
```

```
            name: frontend-service
```

```
            port:
```

```
              number: 80
```

```
      - path: /api       # API path for backend service
```

```
        pathType: Prefix
```

---

backend:

service:

name: backend-service

port:

number: 8080

- path: /static # Static content path (e.g., images, CSS)

pathType: Prefix

backend:

service:

name: static-service

port:

number: 80

---

## Explanation of This Ingress YAML

### 1 Ingress Class & Controller:

- `ingressClassName: nginx` → Specifies that this Ingress should use the Nginx Ingress Controller.

### 2 TLS Configuration:

- The Ingress uses **HTTPS** and references a **TLS secret** (`tls-secret`) for securing the domain `app.example.com`.
- You must create this secret using a **TLS certificate** before applying the Ingress.

### 3 Routing Rules:

- Requests to `app.example.com/` → Routed to **frontend-service** on port **80**.
- Requests to `app.example.com/api` → Routed to **backend-service** on port **8080**.
- Requests to `app.example.com/static` → Routed to **static-service** on port **80**.

### 4 Annotations for Nginx:

- 
- `nginx.ingress.kubernetes.io/rewrite-target: /` → Ensures correct path forwarding.
  - `nginx.ingress.kubernetes.io/ssl-redirect: "true"` → Forces HTTPS redirection.

## Apply the Ingress Configuration

Once the Ingress Controller is running, apply the YAML file using:

```
kubectl apply -f ingress.yaml
```

## Verify the Ingress Resource

Check if the Ingress resource is created successfully:

```
kubectl get ingress
```

You should see an output like:

NAME	CLASS	HOSTS	ADDRESS
PORTS	AGE		

---

my-app-ingress	nginx	app.example.com	192.168.1.100
80,443	5m		

### ✔ Check Ingress Rules

```
kubectl describe ingress my-app-ingress
```

---

# Kubernetes Ingress Example for an E-Commerce

## Platform

Below is a **Kubernetes YAML configuration** that sets up an **NGINX Ingress Controller** to route traffic to different microservices of an e-commerce application.

### Step 1: Deploy Microservices

Let's create Deployments and Services for the **Frontend, Product, Order, and User Services**.

#### Frontend Service (**frontend-service.yaml**)

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: frontend-deployment
```

```
spec:
```

---

```
replicas: 2
```

```
selector:
```

```
  matchLabels:
```

```
    app: frontend
```

```
template:
```

```
  metadata:
```

```
    labels:
```

```
      app: frontend
```

```
spec:
```

```
  containers:
```

```
    - name: frontend
```

```
      image: my-ecommerce-frontend:v1
```

```
      ports:
```

```
        - containerPort: 80
```

---

---

apiVersion: v1

kind: Service

metadata:

name: frontend-service

spec:

selector:

app: frontend

ports:

- protocol: TCP

port: 80

targetPort: 80



---

## Product Service (**product-service.yaml**)

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: product-deployment
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      app: product
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: product
```

```
    spec:
```

---

containers:

- name: product

- image: my-ecommerce-product:v1

- ports:

- containerPort: 8080

---

apiVersion: v1

kind: Service

metadata:

- name: product-service

spec:

- selector:

- app: product

- ports:

---

- protocol: TCP

- port: 8080

- targetPort: 8080

## Order Service (**order-service.yaml**)

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: order-deployment
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      app: order
```

---

template:

metadata:

labels:

app: order

spec:

containers:

- name: order

image: my-ecommerce-order:v1

ports:

- containerPort: 8081

---

apiVersion: v1

kind: Service

metadata:

---

```
  name: order-service
```

```
spec:
```

```
  selector:
```

```
    app: order
```

```
  ports:
```

```
    - protocol: TCP
```

```
      port: 8081
```

```
      targetPort: 8081
```

## User Service (**user-service.yaml**)

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: user-deployment
```

```
spec:
```

---

replicas: 2

selector:

matchLabels:

app: user

template:

metadata:

labels:

app: user

spec:

containers:

- name: user

image: my-ecommerce-user:v1

ports:

- containerPort: 8082

---

```
---

apiVersion: v1

kind: Service

metadata:

  name: user-service

spec:

  selector:

    app: user

  ports:

    - protocol: TCP

      port: 8082

      targetPort: 8082
```

## Step 2: Create an Ingress Resource (**ingress.yaml**)

This Ingress resource defines routing rules based on paths.

---

apiVersion: networking.k8s.io/v1

kind: Ingress

metadata:

name: ecommerce-ingress

annotations:

nginx.ingress.kubernetes.io/rewrite-target: /

spec:

ingressClassName: nginx

rules:

- host: example.com

http:

paths:

- path: /

pathType: Prefix

backend:



---

service:

name: frontend-service

port:

number: 80

- path: /products

pathType: Prefix

backend:

service:

name: product-service

port:

number: 8080

- path: /orders

pathType: Prefix

backend:

service:

---

```
      name: order-service

      port:

        number: 8081

- path: /users

  pathType: Prefix

  backend:

    service:

      name: user-service

      port:

        number: 8082
```

### Step 3: Deploy the NGINX Ingress Controller

If you haven't already installed the **NGINX Ingress Controller**, you can deploy it using Helm:

---

```
helm repo add ingress-nginx
```

```
https://kubernetes.github.io/ingress-nginx
```

```
helm install nginx-ingress ingress-nginx/ingress-nginx
```

#### Step 4: Apply the Configurations

Now, apply all the YAML files to your Kubernetes cluster.

```
kubectl apply -f frontend-service.yaml
```

```
kubectl apply -f product-service.yaml
```

```
kubectl apply -f order-service.yaml
```

```
kubectl apply -f user-service.yaml
```

```
kubectl apply -f ingress.yaml
```

#### Step 5: Access the Services

Once everything is deployed:

- Visit <https://example.com/> → **Frontend Service**
- Visit <https://example.com/products> → **Product Service**

- 
- Visit <https://example.com/orders> → **Order Service**
  - Visit <https://example.com/users> → **User Service**

If you are running this in **minikube**, enable the ingress controller and get the ingress IP:

```
minikube addons enable ingress
```

```
kubectl get ingress
```

### Why Is This a Good Setup?

- ✓ **Single Entry Point:** Users don't need to remember multiple IPs or ports.
- ✓ **Path-Based Routing:** Traffic is routed correctly to the respective microservice.
- ✓ **Cost-Effective:** No need for multiple Load Balancers.
- ✓ **Security & SSL Handling:** TLS termination can be handled at the Ingress level.
- ✓ **Scalability:** Can add more microservices easily.

---

## Host-Based Routing in Kubernetes Ingress

In **host-based routing**, traffic is routed based on the **domain name (host)** rather than the URL path. This is useful when hosting **multiple applications** or **subdomains** on the same Kubernetes cluster.

### Real-Life Example: Hosting Multiple Applications on Different Domains

Imagine you have a **news company** with the following services:

- **Main Website:** `news.example.com`
- **Sports Section:** `sports.example.com`
- **Finance Section:** `finance.example.com`

Instead of creating separate Load Balancers for each, we use a **single Ingress** to route requests based on the hostname.

#### Step 1: Deploy Services

Each service has its own **deployment and service**.

---

## Main Website Service (**news.example.com**)

apiVersion: apps/v1

kind: Deployment

metadata:

name: news-deployment

spec:

replicas: 2

selector:

matchLabels:

app: news

template:

metadata:

labels:

app: news

spec:

---

containers:

- name: news

- image: my-news-app:v1

- ports:

- containerPort: 80

---

apiVersion: v1

kind: Service

metadata:

- name: news-service

spec:

- selector:

- app: news

- ports:

---

- protocol: TCP

- port: 80

- targetPort: 80

## **Sports Section Service (sports.example.com)**

apiVersion: apps/v1

kind: Deployment

metadata:

- name: sports-deployment

spec:

- replicas: 2

- selector:

- matchLabels:

- app: sports

- template:



---

```
  metadata:

    labels:

      app: sports

  spec:

    containers:

      - name: sports

        image: my-sports-app:v1

        ports:

          - containerPort: 80

---

apiVersion: v1

kind: Service

metadata:

  name: sports-service
```

---

spec:

selector:

app: sports

ports:

- protocol: TCP

port: 80

targetPort: 80

## Finance Section Service (**finance.example.com**)

apiVersion: apps/v1

kind: Deployment

metadata:

name: finance-deployment

spec:

replicas: 2

---

selector:

matchLabels:

app: finance

template:

metadata:

labels:

app: finance

spec:

containers:

- name: finance

image: my-finance-app:v1

ports:

- containerPort: 80

---

---

```
apiVersion: v1

kind: Service

metadata:
  name: finance-service

spec:
  selector:
    app: finance

  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

## Step 2: Create an Ingress for Host-Based Routing

```
apiVersion: networking.k8s.io/v1

kind: Ingress
```

---

metadata:

name: news-ingress

annotations:

nginx.ingress.kubernetes.io/rewrite-target: /

spec:

ingressClassName: nginx

rules:

- host: news.example.com

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name: news-service

---

port:

number: 80

- host: sports.example.com

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name: sports-service

port:

number: 80

- host: finance.example.com

http:

paths:

---

- path: /

pathType: Prefix

backend:

service:

name: finance-service

port:

number: 80

### Step 3: Apply Ingress and Update DNS

Apply the ingress:

```
kubectl apply -f ingress.yaml
```

Check the ingress details:

```
kubectl get ingress
```

1. The output will show an **AWS Load Balancer hostname** (if using Kops on AWS).

Add this hostname as **CNAME records** in Route 53 (or any DNS provider)

---

for:

- `news.example.com`
- `sports.example.com`
- `finance.example.com`

### Expected Behavior

- Visiting `https://news.example.com/` → Routes to the **News Service**.
- Visiting `https://sports.example.com/` → Routes to the **Sports Service**.
- Visiting `https://finance.example.com/` → Routes to the **Finance Service**.

Each service remains independent while being managed under a **single Ingress and Load Balancer**.

### Why Use Host-Based Routing?

- ✓ **Efficient:** No need for multiple Load Balancers.
- ✓ **Cost-Effective:** Reduces AWS ELB expenses.



---

✓ **Scalable:** Easily add new subdomains in the Ingress rule.

✓ **Simplifies Management:** Centralized control over multiple services.



# Complete Example Scenario of Using Ingress in Kubernetes (with kOps)



## Scenario:

You are running a **Kubernetes cluster created using kOps** and want to expose two different applications (**frontend** and **backend**) using a **single domain** (**example.com**) with **Ingress for routing traffic**.



## What We Will Cover in This Example



**Host-based Routing:** Route **api.example.com** to backend service & **app.example.com** to frontend service.



**Path-based Routing:** Serve different services based on **/api** and **/app**.



**TLS/HTTPS Support:** Secure traffic using an SSL certificate.



**Nginx Ingress Controller Setup:** Install Ingress in a kOps cluster.



**Complete Deployment, Service & Ingress YAMLs:** To make it fully working.

---

## ◆ Step 1: Install an Ingress Controller on Your kOps Cluster

Since kOps does not come with an Ingress Controller by default, you need to install one:

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/kubernetes/ingress-nginx/main/deploy/static/provider/cloud/deploy.yaml
```

This installs **NGINX Ingress Controller**, which will manage traffic routing.

## ◆ Step 2: Create Deployments and Services for Frontend & Backend

We need two applications:

- Frontend (**frontend-service**) → Runs on port **80**.
- Backend (**backend-service**) → Runs on port **5000**.

📌 Deployment & Service for Frontend (React, Angular, or Vue App)

---

apiVersion: apps/v1

kind: Deployment

metadata:

name: frontend

labels:

app: frontend

spec:

replicas: 2

selector:

matchLabels:

app: frontend

template:

metadata:

labels:

app: frontend

---

```
spec:

  containers:

    - name: frontend

      image: myfrontend:latest

      ports:

        - containerPort: 80

---

apiVersion: v1

kind: Service

metadata:

  name: frontend-service

spec:

  selector:

    app: frontend

  ports:
```

---

- protocol: TCP

port: 80

targetPort: 80

type: ClusterIP

## Deployment & Service for Backend (Node.js, Django, or Flask API)

apiVersion: apps/v1

kind: Deployment

metadata:

name: backend

labels:

app: backend

spec:

replicas: 2

selector:

---

```
  matchLabels:

    app: backend

template:

  metadata:

    labels:

      app: backend

  spec:

    containers:

      - name: backend

        image: mybackend:latest

        ports:

          - containerPort: 5000

---

apiVersion: v1

kind: Service
```

---

metadata:

name: backend-service

spec:

selector:

app: backend

ports:

- protocol: TCP

port: 5000

targetPort: 5000

type: ClusterIP

### ◆ Step 3: Create an Ingress Resource for Routing Traffic

Now, we define an **Ingress resource** that:

- ◆ Routes **app.example.com** to frontend-service
- ◆ Routes **api.example.com** to backend-service
- ◆ Routes **example.com/api** to backend-service



- 
- ◆ Routes **example.com/app** to frontend-service
  - ◆ Enables **HTTPS/TLS termination**

### Complete Ingress YAML

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: my-ingress
```

```
  annotations:
```

```
    nginx.ingress.kubernetes.io/rewrite-target: /
```

```
    nginx.ingress.kubernetes.io/ssl-redirect: "true"
```

```
spec:
```

```
  ingressClassName: nginx
```

```
  tls:
```

```
  - hosts:
```

```
    - example.com
```

---

- app.example.com

- api.example.com

secretName: tls-secret

rules:

- host: app.example.com

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name: frontend-service

port:

number: 80

- host: api.example.com

---

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name: backend-service

port:

number: 5000

- host: example.com

http:

paths:

- path: /app

pathType: Prefix

backend:

---

```
    service:

      name: frontend-service

      port:

        number: 80

  - path: /api

    pathType: Prefix

    backend:

      service:

        name: backend-service

        port:

          number: 5000
```

#### ◆ Step 4: Configure TLS for Secure HTTPS

To enable HTTPS for your Ingress, you need to create a TLS secret.

##### Generate TLS Certificate (Self-Signed)

---

```
openssl req -x509 -nodes -days 365 -newkey rsa:2048 -keyout  
tls.key -out tls.crt -subj "/CN=example.com/O=example.com"
```

### Create Kubernetes Secret

```
kubectl create secret tls tls-secret --cert=tls.crt  
--key=tls.key
```

## ◆ Step 5: Apply Everything to the Cluster

Now, apply the deployments, services, and ingress:

```
kubectl apply -f frontend.yaml
```

```
kubectl apply -f backend.yaml
```

```
kubectl apply -f ingress.yaml
```

## ◆ Step 6: Test the Ingress Setup

Once the Ingress is deployed, test it using `curl`:

---

```
curl -k https://example.com/api # Should return backend  
response
```

```
curl -k https://example.com/app # Should return frontend  
response
```

```
curl -k https://app.example.com # Should return frontend  
response
```

```
curl -k https://api.example.com # Should return backend  
response
```

Or, open a **web browser** and access:

- ◆ <https://app.example.com> (Frontend)
- ◆ <https://api.example.com> (Backend)
- ◆ <https://example.com/api> (Backend API)
- ◆ <https://example.com/app> (Frontend UI)

---

## Final Summary

- ✓ KOps Cluster + Ingress Controller Installed
- ✓ Frontend & Backend Services Deployed
- ✓ Host-based & Path-based Routing Configured
- ✓ TLS/HTTPS Secured
- ✓ Ingress Successfully Implemented & Tested

With this setup, your **Kubernetes applications are now securely exposed to the internet using Ingress!** 🎉

---

## ◆ In Summary: Mastering Kubernetes Ingress

Kubernetes Ingress is a **powerful way to expose and manage external access** to your services efficiently. Before Ingress, we relied on NodePort and LoadBalancer services, which were either inflexible or expensive. **Ingress solves these challenges** by providing a centralized way to route traffic, apply security policies, and enable load balancing—all within Kubernetes.

In this example, we covered:

- ✓ **Installing an Ingress Controller** in a Kubernetes cluster (using kOps).
- ✓ **Deploying and Exposing Services** (Frontend & Backend) inside the cluster.
- ✓ **Using Host-Based and Path-Based Routing** to direct traffic efficiently.
- ✓ **Enabling TLS/HTTPS for Secure Access** using a self-signed certificate.
- ✓ **Testing & Validating Ingress Rules** to ensure smooth traffic flow.

With Ingress, you **simplify traffic management, improve security, and optimize resource usage**, making it an essential tool for production-grade Kubernetes environments! 🚀



---

## ◆ Wrapping Up: Thank You for Being Part of This

### Journey!

And that's a wrap on **Kubernetes Ingress!** 🎉 We've explored everything from why Ingress is needed, how it works, key features, configuration, and real-world scenarios to implement it effectively. **Ingress is a powerful tool** that simplifies routing, enhances security, and optimizes traffic management in Kubernetes clusters. I hope this guide gave you a **clear and practical** understanding of how to leverage Ingress in your Kubernetes environment. 🚀

But this is just one step in our **Kubernetes learning journey!** 💡 There's so much more to uncover, and I'm excited to continue sharing **daily hands-on tasks, practical insights, and deep dives into key Kubernetes concepts** with all of you.

🙏 **A huge thank you for supporting me throughout this journey!** Your engagement, encouragement, and enthusiasm keep me motivated to bring more valuable content every day. Your likes, comments, and shares make this series even more special! ❤️

---

This journey doesn't stop here—**there's more to come!** 🔥 If you've been finding these posts helpful, **make sure to follow me** to stay updated with more **Kubernetes deep dives, best practices, and real-world hands-on tasks!**

Let's keep learning, growing, and building together. See you in the next one! 🚀 ✨

**Stay tuned, stay curious, and keep exploring Kubernetes!** 🔥💡